

An interleaving approach to combinatorial testing and failure-inducing interaction identification^{*}

Xintao Niu and Changhai Nie
State Key Laboratory for Novel
Software Technology
Nanjing University
China, 210023
changhainie@nju.edu.cn

Hareton Leung
Department of computing
Hong Kong Polytechnic
University
Kowloon, Hong Kong
cshleung@comp.polyu.edu.hk

Jeff Lei
Department of Computer
Science and Engineering
The University of Texas at
Arlington
ylei@cse.uta.edu

JiaXi Xu and Yan Wang
School of Information
Engineering
Nanjing Xiaozhuang University
China, 211171
xujiaxi@njzuc.edu.cn

Xiaoyin Wang
Department of Computer
Science
University of Texas at San
Antonio
Xiaoyin.Wang@utsa.edu

ABSTRACT

Combinatorial testing(CT) seeks to detect potential faults caused by various interactions of factors that can influence the software systems. When applying CT, it is a common practice to first generate a set of test cases to cover each possible interaction and then to identify the failure-inducing interaction after a failure is detected. Although this conventional procedure is simple and straightforward, we conjecture that it is not the ideal choice in practice. This is because 1) testers desire to identify the root cause of failures before all the needed test cases are generated and executed 2) the early identified failure-inducing interactions can guide the remaining test cases generation, such that many unnecessary and invalid test cases can be avoided. For these reasons, we propose a novel CT framework that allows both generation and identification process to interact with each other. As a result, both generation and identification stages will be done more effectively and efficiently. We conducted a series of empirical studies on several open-source software, the results of which show that our framework can identify the failure-inducing interactions more quickly than traditional approaches, while requiring fewer test cases.

^{*}This work was supported by the National Natural Science Foundation of China (No. 61272079), the Research Fund for the Doctoral Program of Higher Education of China (No.20130091110032), the Science Fund for Creative Research Groups of the National Natural Science Foundation of China(No. 61321491), and the Major Program of National Natural Science Foundation of China (No. 91318301)

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

WOODSTOCK '97 El Paso, Texas USA

Copyright 20XX ACM X-XXXXX-XX-X/XX/XX ...\$15.00.

Categories and Subject Descriptors

D.2.5 [Software Engineering]: Testing and debugging—*Debugging aids, testing tools*

General Terms

Reliability, Verification

Keywords

Software Testing, Combinatorial Testing, Covering Array, Failure-inducing interactions

1. INTRODUCTION

Modern software is becoming more and more complex. To test such software is challenging, as the candidate factors that can influence the system's behaviour, e.g., configuration options, system inputs, message events, are enormous. Even worse, the interactions between these factors can also crash the system, e.g., the incompatibility problems. In consideration of the scale of the real industrial software, to test all the possible interactions of all the factors (we call them the interaction space) is not feasible, and even if it is possible, it is not wise to test all the interactions because most of them do not provide any useful information.

Many empirical studies show that, in real software systems, the effective interaction space, i.e., targeting fault detection, makes up only a small proportion of the overall interaction space [20, 21]. What's more, the number of factors involved in these effective interactions is relatively small, of which 4 to 6 is usually the upper bounds[20]. With this observation, applying CT in practice is appealing, as it is proven to be effective to detect the interaction faults in the system.

A typical CT life-cycle is shown in Figure 1, which contains four main testing stages. At the very beginning of the testing, engineers should extract the specific model of the software under test. In detail, they should identify the factors, such as user inputs, configure options, that could affect the system's behavior. Further effort is needed to figure out

the constraints and dependencies between each factor and corresponding values for valid testing. After the modeling stage, a set of test cases should be generated and executed to expose the potential faults in the system. In CT, each test case is a set of assignments of all the factors in the test model. Thus, when such a test case is executed, all the interactions contained in the test case are deemed to be checked. The main target of this stage is to design a relatively small set of test cases to achieve some specific coverage. The third testing stage in this cycle is the fault localization, which is responsible for identifying the failure-inducing interactions. To characterize the failure-inducing interactions of corresponding factors and values is important for future bug fixing, as it will reduce the suspicious code scope to be inspected. The last testing stage of CT is evaluation. In this stage, testers will assess the quality of the previously conducted testing tasks. If the assessment result shows that the previous testing process does not fulfil the testing requirement, some testing stages should be improved, and sometimes, may even need to be re-conducted.

Although this conventional CT framework is simple and straightforward, in terms of the test cases generation and fault localization stages, we conjecture that first-generation-then-identification is not the proper choice in practice. The reasons are twofold. First, it is not realistic for testers to wait for all the needed test cases are generated before they can diagnose and fix the failures that have been detected [37]; Second, and the most important, utilizing the early determined failure-inducing interactions can guide the following test cases generations, such that many unnecessary and invalid test cases can be avoided. For this we get the key idea of this paper: *Generation and Fault Localization process should be integrated*.

Based on the idea, we propose a new CT framework, which instead of dividing the generation and identification into two independent stages, it integrates these two stages into one. Specifically, we first execute one or more tests until a failure is observed. Next we immediately turn to the fault localization stage, i.e., identify failure-inducing interactions for that failure. These failure-inducing interactions are used to update the currently coverage. In particular, interactions that are related to these failure-inducing interactions do not need to be covered in future executions. Then, we go back to perform regular combinatorial testing and continue.

We remodel the test cases generation and failure-inducing interactions identification modules to make them better adapt to this new framework. Specifically, for the generation part of our framework, we augment it by forbidding the appearance of test cases which contain the identified failure-inducing interactions. This is because those test cases contain a failure-inducing interaction will fail as expected, so that it makes no sense for the further failure detection. For the failure-inducing identification module, we augment it by making them to contribute more coverage. More specifically, we refine the additional test cases generation in this module, so that it can not only help to identify the failure-inducing interactions, but also cover as many uncovered interactions as possible. As a result, our new CT framework needs fewer test cases than traditional CT.

We conducted a series of empirical studies on several open-source software to evaluate our new framework. Specifically, we performed two comparisons studies. The first one is to compare our new framework with the traditional one, which



Figure 1: The life cycle of CT

first generates a complete set of test cases and then performs the fault localization. The second one is to compare our framework with the Feedback-driven CT [11, 36], which also adapts an iterative framework to generate test cases and identifying failure-inducing interactions, but to address the problem of inadequate testing. The results show that, in terms of test case generation and failure-inducing interactions identification, our approach can significantly reduce the overall needed test cases and as a result it can more quickly identify the failure-inducing interactions of the system under test.

The main contributions of this paper are as follows.

1. We propose a new CT framework which combines the test cases generation and fault localization more closely.
2. We augment the traditional CT test cases generation and failure-inducing interactions identification process to make them adapt to the new framework.
3. We perform a series of comparisons with traditional CT and Feedback-driven CT. The result of the empirical studies are discussed.

The rest of the paper is organised as follows: Section 2 presents the preliminary background of CT. Section 3 presents an motivating example. Section 4 describes our new framework and a simple case study is also given. Section 5 gives the limitation of our approach as well as some corresponding measures to handle it. Section 6 describes the experiment subjects and its input model. Section 7 presents the empirical studies and the results are discussed. Section 8 shows the related works. Section 9 concludes the paper and proposes some further work.

2. BACKGROUND

This section presents some definitions and propositions to give a formal model for CT.

Assume that the Software Under Test (SUT) is influenced by n parameters, and each parameter p_i can take the values from the finite set V_i ($i = 1, 2, \dots, n$). The definitions below are originally defined in [27].

Definition 1. A test case of the SUT is a tuple of n values, one for each parameter of the SUT. It is denoted as (v_1, v_2, \dots, v_n) , where $v_1 \in V_1, v_2 \in V_2 \dots v_n \in V_n$.

In practice, these parameters in the test case can represent many factors, such as input variables, run-time options, building options or various combination of them. We need to execute the SUT with these test cases to ensure the correctness of the behaviour of the SUT.

We consider any abnormally executing test case as a *fault*. It can be a thrown exception, compilation error, assertion failure or constraint violation. When faults are triggered by some test cases, it is desired to figure out the cause of these faults. Some subsets of these test cases should be analysed.

Definition 2. For the SUT, the n -tuple $(-, v_{n_1}, \dots, v_{n_k}, \dots)$ is called a k -degree *schema* ($0 < k \leq n$) when some k parameters have fixed values and other irrelevant parameters are represented as ”-”.

In effect a test case itself is a k -degree *schema*, when $k = n$. Furthermore, if a test case contains a *schema*, i.e., every fixed value in the schema is in this test case, we say this test case *contains* the *schema*.

Note that the schema is a formal description of the interaction between parameter values we discussed before.

Definition 3. Let c_l be a l -degree schema, c_m be an m -degree schema in SUT and $l < m$. If all the fixed parameter values in c_l are also in c_m , then c_m *subsumes* c_l . In this case we can also say that c_l is a *sub-schema* of c_m and c_m is a *super-schema* of c_l , which can be denoted as $c_l \prec c_m$.

For example, the 2-degree schema $(-, 4, 4, -)$ is a sub-schema of the 3-degree schema $(-, 4, 4, 5)$, that is, $(-, 4, 4, -) \prec (-, 4, 4, 5)$.

Definition 4. If all test cases that contain a schema, say c , trigger a particular fault, say F , then we call this schema c the *faulty schema* for F . Additionally, if none of sub-schema of c is the *faulty schema* for F , we then call the schema c the *minimal failure-causing schema (MFS)* [27] for F .

Note that MFS is identical to the failure-inducing interaction we discussed previously. In this paper, the terms *failure-inducing interactions* and *MFS* are used interchangeably. Figuring the MFS out helps to identify the root cause of a failure and thus facilitate the debugging process.

2.1 CT Test Case Generation

When applying CT, the most important work is to determine whether the SUT suffers from the interaction faults or not, i.e., to detect the existence of the MFS. Rather than impractically executing exhaustive test cases, CT commonly design a relatively small set of test cases to cover all the schemas with the degree no more than a prior fixed number, t . Such a set of test cases is called the *covering array*. If some test cases in the covering array failed in execution, then the interaction faults is regard to be detected.

Many studies in CT focus on how to generate such a test suite with the aim of making the size of the test suite as small as possible. In general, most of these studies can be classified into three categories according to the construction strategy of the covering array:

1) One test case one time : This strategy repeats generating one test case as one row of the covering array and counting the covered schemas achieved until all schemas are covered.

2) A set of test cases one time: This strategy generates a set of test cases at each iteration. Through mutating the values of some parameters of some test cases in this test set, it focuses optimising the coverage. If the coverage is finally satisfied, it will reduce the size of the set to see if fewer test

cases can still fulfil the coverage. Otherwise, it will increase the size of test set to cover all the schemas[5].

3) Others : This strategy differentiates from the previous two strategies in that it does not firstly generate complete test cases [23]. Instead, it first focuses on assigning values to some part of the factors or parameters to cover the schemas that related to these factors, and then fills up the remaining part to form complete test cases.

In this paper, we focus on the first strategy: One test case one time as it immediately get a complete test case such that the testers can execute and diagnose in the early stage. As we will see later, with respect to the MFS identification, this strategy is the most flexible and efficient one comparing with the other two strategies.

2.2 Identify the failure-inducing interactions

To detect the existence of MFS in the SUT is still far from figuring out the root cause of the failure, as we do not know exactly which schemas in the failed test cases should be responsible for the failure. In fact, for a failing test case (v_1, v_2, \dots, v_n) , there can be at most $2^n - 1$ possible schemas for the MFS. Hence, more test cases should be generated to identify the MFS. In CT, the main work in fault localization is to identify the failure-inducing interactions. So in this paper we only focus on the MFS identification. Further works of fault localization such as isolating the defect inside the source code will not be discussed.

A typical MFS identification process is shown in Table 1. This example assumes the SUT has 3 parameters, each can take 2 values, and the test case $(1, 1, 1)$ fails. Then in Table 1, as test case t failed, we mutate one factor of the t one time to generate new test cases: $t_1 - t_3$. It turns out that test case t_1 passed, which indicates that this test case break the MFS in the original test case t . So $(1, -, -)$ should be a failure-causing factor, and as other mutating process all failed, which means no other failure-inducing factors were broken, therefore, the MFS in t is $(1, -, -)$.

Table 1: OFOT example

	Original test case			Outcome
t	1	1	1	Fail
Additional test cases				
t_1	0	1	1	Pass
t_2	1	0	1	Fail
t_3	1	1	0	Fail

This identification process mutate one factor of the original test case at a time to generate extra test cases. And then according to the outcome of the test cases execution result, it will identify the MFS of the original failing test cases. It is called the OFOT method, which is a well-known MFS identification method in CT. In this paper, we will focus on this identification method. It should be noted that the following proposed CT framework can be easily applied to other MFS identification methods.

3. MOTIVATING EXAMPLE

In this section, a motivating example is presented to show how traditional CT works as well as its limitations. This example is derived from our attempt to test a real-world software-HSQLDB, which is a pure-java relational database

engine with large and complex configuration space. To extract and manipulate valid configurations of this highly-configurable system is important, as different configuration can result in significant different behaviours of the system [16, 31, 33] (HSQLDB may normally works under some properly configurations, but crashes or throws exceptions under some other configurations, for example).

Considering the large configuration space of HSQLDB, we firstly utilized CT to generate a relatively small set of test cases. Each of them is actually a set of specific assignments to those options we cared. For each configuration, HSQLDB is tested by sending prepared SQL commands. We recorded the output of each run, but unfortunately, about half of them produced exceptions or warnings. Following the schedule of traditional CT, we started the identification process to isolate the failure-inducing option interactions in those failing configurations. Each failing configuration should be individually handled, in principle, as there may exist distinct failure-inducing option interactions among them. However, this successively identification process, although appealing, was hardly ever followed for this case study. This is because there are too many failing configurations and most of them contain the same failure-inducing option interactions, based on which the MFS identification process is wasteful and inefficient.

For the sake of convenience, we provide a highly simplified scenario to illustrate the problems we encountered. Consider four options in HSQLDB – *Server type*, *Scroll Type*, *Parameterised SQL* and *Statement Type*. The possible values each option can take on are shown in Table 2. Based on the report in the bug tracker of HSQLDB¹, an *incompatible exception* will be triggered if a *parameterised sql* is executed as *prepared statement* by HSQLDB. Hence, when option *parameterisedSQL* is set to be *true* and *Statmetent* to be *preparedStatement*, our testing will crash. Besides this failure, there exists another option value which can also crash this database engine. It is when *Scroll Type* is assigned to *sensitive*, as this feature is not supported by this version of HSQLDB². Without these knowledge at prior, we needed to detect and isolate these two failure-inducing option interactions by CT.

Table 2: Highly simplified configuration of HSQLDB

Option		Values
o_1	Server type	server, web-server, in-process
o_2	Scroll type	sensitive,insensitive, forward-only
o_3	parameterised SQL	true, false
o_4	Statement Type	statement, preparedStatement

Table 3 illustrates the process of traditional CT on this subject. For simplicity of notation, we use consecutive symbols 0, 1, 2 to represent each value of each option (For *parameterisedSQL* and *Statement*, the symbol is up to 1). According to Table 3, traditional CT first generated and executed the 2-way covering array ($t_1 - t_9$ in the *generation* part). Note that this covering array covered all the 2-degree schemas for the SUT. For example, all the possible values combinations between first and second parameters appear

¹details see: <http://sourceforge.net/p/hsqldb/bugs/1173/>

²details see: <http://hsqldb.org/doc/guide/guide.html>

at least one test case.

After detecting that t_1 , t_4 , and t_7 failed during testing, it is desired to respectively identify the MFS of these failing test cases. For t_1 , OFOT method is used to generate four additional test cases ($t_{10} - t_{13}$), and the MFS $(-, 0, -, -)$ of t_1 is identified (*Scroll Type* is assigned to *sensitive*, respectively). This is because only when changing the second factor of t_1 , the extra-generated test case will pass. Then the same process is applied on t_4 and t_7 . Finally, we found that the MFS of t_4 is $(-, -, -, -)$, indicating that OFOT failed to determine the MFS (will be discussed later), and the MFS of t_7 is the same as t_1 . Totally, for detecting and identifying the MFS in this example, we have generated 12 additional test cases (marked with star).

Table 3: Traditional generation-identification life-cycle

Generation					
test case					Outcome
	o_1	o_2	o_3	o_4	
t_1	0	0	0	0	Fail
t_2	0	1	1	1	Pass
t_3	0	2	1	0	Pass
t_4	1	0	0	1	Fail
t_5	1	1	0	0	Pass
t_6	1	2	1	1	Fail
t_7	2	0	1	1	Pass
t_8	2	1	0	0	Pass
t_9	2	2	0	0	Fail

Identification						
t_1 (0,0,0,0)	t_{10}^*	1	0	0	0	Fail
	t_{11}^*	0	1	0	0	Pass
	t_{12}^*	0	0	1	0	Fail
	t_{13}^*	0	0	0	1	Fail
MFS		(-, 0, -, -)				
t_4 (1,0,0,1)	t_{14}^*	2	0	0	1	Fail
	t_{15}^*	1	1	0	1	Fail
	t_{16}^*	1	0	1	1	Fail
	t_{17}^*	1	0	0	0	Fail
MFS		(-, -, -, -)				
t_7 (2,0,1,1)	t_{18}^*	0	0	1	1	Fail
	t_{19}^*	2	1	1	1	Pass
	t_{20}^*	2	0	0	1	Fail
	t_{21}^*	2	0	1	0	Fail
MFS		(-, 0, -, -)				

We refer to such traditional life-cycle as **Sequential CT** (SCT). However, we believe this may not be the best choice in practice. The first reason is that the engineers normally do not want to wait for fault localization after all the test cases are executed. The early bug fixing is appealing and can give the engineers confidence to keep on improving the quality of the software. The second reason, which is also the most important, is such life-cycle can generate many redundant and unnecessary test cases, which negatively impacted on both test cases generation and MFS identification. The most obvious negative effect in this example is that we did not identify the expected failure-inducing interaction $(-, -, 0, 1)$, which corresponds that option *parameterisedSQL* is set to be *true* and *Statmetent* to be *preparedStatement*. More shortcomings of the sequential CT are discussed as

following:

3.1 Redundant test cases

The first shortcoming of SCT is that it may generate redundant test cases, such that some of them do not cover as more uncovered schemas as possible. As a consequent, SCT may generate more test cases than actually needed. This can be reflected in the following two aspects:

1) The test cases generated in the identification stage can also contribute some coverage, i.e., the schemas appear in the passing test cases in the identification stage may have already been covered in the test cases generation stage. For example, when we identify the MFS of t_1 in Table 3, the schema $(0, 1, -, -)$ contained in the extra passing test case $t_{11} - (0, 1, 0, 0)$ has already been appeared in the passing test case $t_2 - (0, 1, 1, 1)$. In other word, if we firstly identify the MFS of t_1 , then t_2 is not a good choice as it does not covered as many 2-degree schemas as possible. For example, $(1, 1, 1, 1)$ is better than this test case at contributing more coverage.

2) The identified MFS should not appear in the following generated test cases. This is because according to the definition of MFS, each test case containing this schema will trigger a failure, i.e., to generate and execute more than one test case contained the MFS makes no sense for the failure detection. Take the example in Table 3, after identifying the MFS $(-, 0, -, -)$ of t_1 , we should not generate the test case t_4 and t_7 . This is because they also contain the identified MFS $(-, 0, -, -)$, which will result in them failing as expected. Since the expected failure caused by MFS $(-, 0, -, -)$ makes t_7 and t_9 superfluous for error-detection, the additional test cases (t_{14} to t_{21}) generated for identifying the MFS in t_4 and t_7 are also not necessary.

3.2 Multiple MFS in one test case

Multiple MFS will negatively affect the accurateness of MFS identification – and even make it hardly obtain a validate schema. For example, there are two MFS in t_4 in Table 3, i.e., $(-, 0, -, -)$ and $(-, -, 0, 1)$ (in bold). When we use OFOT method, we found all the extra-generated test cases (t_{14} to t_{17}) failed. These outcomes give OFOT a false indication that all the failure-inducing factors are not broken by mutating those four parameter values. As a result, OFOT cannot determine which schemas are MFS, which is denoted as $(-, -, -, -)$.

The reason why OFOT cannot properly work is that this approach can only break one MFS at a time. If there are multiple MFS in one test case, the extra-generated test cases will always fail as they contain other non-broken MFS (see bold parts of t_{14} to t_{17}). Some approaches are proposed to handle this problem, but either can not handle multiple MFS that have overlapping parts [38], or consumes too many extra-generated test cases [34, 30]. So in practice, to make MFS identification more effective and efficient, we need to avoid the appearance of multiple MFS in one test case.

SCT, however, does not offer much support for this concern. This is mainly because it is essentially a post-analysis framework, i.e., the analysis for MFS comes after the completion of test cases generation and execution. As a result, in the generation stage, testers have no knowledge of the possible MFS, and surely it has opportunities that multiple MFS appear in one test case.

3.3 Masking effects

Traditional covering array usually offer an inadequate testing due to *Masking effects* [11, 36]. A masking effect [36] is an effect that some failures or exceptions prevent a test case from testing all valid schemas in that test case, which the test case is normally expected to test. For example in Table 3, t_1 is initially expected to cover all the 2-degree schemas $(0, 0, -, -)$, $(0, -, 0, -)$, $(0, -, -, 0)$, $(-, 0, 0, -)$, $(-, 0, -, 0)$, $(-, -, 0, 0)$, but when we found that it failed during testing, this conclusion is no longer that sure. This is because, the failing of t_1 (*ScrollType* is set to be *sensitive*) crashed HSQLDB, and as a result, it did not go on executing the remaining test code, which may affect the examination of some interactions of t_1 . Hence, we cannot ensure we thoroughly exercise all the interactions in this failing test case.

Since traditional covering array cannot reach a adequate testing, as an alternative, *tested t-way interaction criterion* as a more rigorous coverage standard is proposed [36]. This criterion, to be simply, shows that a t -degree schema is covered iff (1) it appears in a passing test case (2) it is identified as MFS or faulty schema. Now it is surely that this criteria is not satisfied with traditional covering array alone, next let us examine whether this criterion can be satisfied with SCT, i.e., the combination of traditional covering array and MFS identification.

One obvious insight is that if there is only **single** MFS in each failing test case, this criterion is satisfied. This conclusion is based on the fact that the MFS identification is actually a process to isolate the failure-inducing interaction among other interactions in the failing test case, and since there is only single MFS, then other schemas can be determined as non-MFS.

For example in Table 3, t_1 contained a single MFS $(-, 0, -, -)$, and we identified this MFS by generating four extra test cases (t_{10} to t_{13}). As for t_1 , the schema $(-, 0, -, -)$ is determined to be MFS, but since the target of that testing is 2-way coverage, i.e., to cover all the 2-degree schemas, this 1-degree schema do not contribute any more coverage. Based on the fact $(-, 0, -, -)$ is MFS, all the test cases contain this schema will fail by definition, and surely the super-schemas of $(-, 0, -, -)$ in this test case $(0, 0, -, -)$, $(-, 0, 0, -)$ and $(-, 0, -, 0)$ are also faulty schemas as all the test cases contain these schemas must contain the MFS $(-, 0, -, -)$, which will fail after execution. The remaining 2-degree schemas $(0, -, 0, -)$, $(0, -, -, 0)$, $(-, -, 0, 0)$ are contained in the extra-generated test case t_{11} $(0, 1, 0, 0)$ (Note that for single MFS, there will be at least one passing extra-generated test case), which are of course non-faulty schemas. After all, all the 2-degree schemas in the failing test case t_1 are satisfied the *tested t-way interaction criterion*.

When problem comes to a failing test case with **multiple** MFS, however, SCT fails not meet that criterion. As discussed previously, SCT cannot properly work on test cases with multiple MFS – and even cannot obtain a validate schema. With this in mind, we cannot determine which schemas in this failing test case are MFS or not, consequently, we cannot ensure we have examined all the t -degree schemas in this failing test case. For example, t_4 has two MFS $(-, 0, -, -)$, $(-, -, 0, 1)$, which did not been identified with OFOT approach (In fact, there is no passing extra-generated test case). As a result, there are two 2-degree schemas $(1, 0, -, -)$ $(-, -, 0, 1)$ in this test case that, are neither contained in a passing test case nor determined as MFS



Figure 2: The Interleaving Framework

or faulty schemas. Hence, *tested t -way interaction criterion* is not satisfied. Since multiple MFS in a test case can introduce masking effects, SCT must be negatively affected by this effect as it lacks mechanisms to avoid the appearance of multiple MFS in failing test cases.

4. INTERLEAVING APPROACH

To overcome such deficiencies in traditional CT, we propose a new CT generation-identification framework – *Interleaving CT* (ICT). Our new framework aims at enhancing the interaction of generation and identification to reduce the unnecessary and invalid test cases discussed previously.

4.1 Overall framework

The basic outline of our framework is illustrated in Figure 2. Specifically, this new framework works as follows: First, it checks whether all the needed schemas are covered or not. Commonly the target of CT is to cover all the t -degree schemas, with t normally be assigned as 2 or 3. Then if the current coverage is not satisfied, it will generate a new test case to cover as many uncovered schemas as possible. After that, it will execute this test case with the outcome of the execution either be pass (executed normally, i.e., does not triggered an exception, violate the expected oracle or the like) or fail (on the contrary). When the test case passes the execution, we will recompute the coverage state, as all the schemas in the passing test case are regarded as error-irrelevant. As a result, the schemas that was not covered before will be determined to be covered if it is contained in this newly generated test case. Otherwise if the test case fails, then we will start the MFS identification module to identify

the MFS in this failing test case. One point that needs to be noted is that if the test case fails, we will not directly change the coverage, as we can not figure out which schemas are responsible for this failure among all the schemas in this test case until we identify them.

The identification module works in a similar way as traditional independent MFS identify process, i.e., repeats generating and executing additional test cases until it can get enough information to diagnose the MFS in the original failing test case. The difference from traditional MFS identifying process is that we record the coverage that this module has contributed to the overall coverage. In detail, when the additional test case passes, we will label the schemas in these test cases as covered if it has not been covered before. And when the MFS is found at the end of this module, we will first set them as forbidden schemas that latter generated test cases should not contain (Otherwise the test case must fail and it cannot contribute to more coverage), and second all the t -degree schemas that are *related* to these MFS as covered. Here the *related* schemas indicate the following three types of t -degree schemas:

First, the MFS **themselves**. Note that we do not change the coverage state after the generated test case fails (both for the generation and identification module), so these MFS will never be covered as they always appear in these failing test cases.

Second, the schemas that are the **super-schemas** of these MFS. By definition of the super-schemas (Definition 3), if the test case contains the super-schemas, it must also contain all its sub-schemas. So every test case that contains the super-schemas of the MFS must fail after execution. As a result, they will never be covered as we do not change the coverage state for failing test cases.

Third, those **implicit forbidden** schemas, which was first introduced in [7]. This type of schemas are caused by the conjunction of multiple MFS. For example, for a SUT with three parameters, and each parameter has two values, i.e., SUT(2, 2, 2). If there are two MFS for this SUT, which are (1, -, 1) and (0, 1, -). Then the schema (-, 1, 1) is the implicit forbidden schema. This is because for any test case that contain this schema, it must contain either (1, -, 1) or (0, 1, -). As a result, (-, 1, 1) will never be covered as all the test cases containing this schema will fail and so we will not change the coverage state. In fact, by Definition 4, they can be deemed as faulty schemas.

As we all know, the terminating condition of the CT framework is to cover all the t -degree schemas. Then since the three types of schemas will never be covered in our new CT framework, we can set them as covered after the execution of the identification module, so that the overall process can stop.

4.2 Modifications of CT activities

More details of the modifications of CT activities are as follows:

(1) *Modified CT Generation*: We adopt the *one test case one time* method as the basic skeleton of the generation process. Originally, the generation of one test case can be formulated as EQ1.

$$t \leftarrow \text{select}(\mathcal{T}_{all}, \Omega, \xi) \quad (\text{EQ1})$$

There are three factors that determine the selection of test case t . \mathcal{T}_{all} represents all the valid test cases that can be se-

lected to be executed. Usually the test cases that have been tested will not be included as they have no more contribution to the coverage. Ω indicates the set of schemas that have not been covered yet. ξ is a random factor. Most CT generation approaches prefer to select a test case that can cover as many uncovered schemas as possible. This greedy selection process does not guarantee an optimal solution, i.e., the final size of the set of test cases is not guaranteed to be minimal. The random factor ξ is used to help to escape from local optimum.

As discussed in Section 3, we should make the MFS not appear in the test cases generated afterwards, by treating them as the forbidden schemas. In other words, the candidate test cases that can be selected are reduced, because those test cases that contain the currently identified MFS should not appear next. Formally, let \mathcal{T}_{MFS} indicates the set of test cases that contain the currently identified MFS, then the test case selection is augmented as EQ2.

$$t \leftarrow \text{select}(\mathcal{T}_{all} - \mathcal{T}_{MFS}, \Omega, \xi) \quad (\text{EQ2})$$

In this formula, the only difference from EQ1 is that the candidate test cases that can be selected are changed to $\mathcal{T}_{all} - \mathcal{T}_{MFS}$, which excludes \mathcal{T}_{MFS} from candidate test cases.

In practice, it is impossible to thoroughly search the exhaustive candidate test cases \mathcal{T}_{all} to get a specific test case t . Hence, some heuristic methods are used to simplify the selection. For example, AETG [4] successively assigns the value to each parameter to form a test case (in random order). The value to be assigned is the one that appears in the greatest number of uncovered schemas. Correspondingly, to exclude the test cases in \mathcal{T}_{MFS} , it is common to utilize a constraint solver to avoid the forbidden schemas [7, 8].

(2) *Modified identification of MFS* : Traditional MFS identification aims at finding the MFS in a failing test case. As discussed before, test cases in the covering array are not enough to identify the MFS. Hence, additional test cases should be generated and executed in this CT activity. Generally, an additional test case is generated based on the original failing test case, so that the failure-inducing parts can be determined by comparing the difference between the additional test cases and the original failing test cases. Take the OFOT approach as an example. In Table 3, the additional test case t_{11} is constructed by mutating the second parameter value of the original failing test case t_1 . Then as t_{11} passed the testing, we can determine that the second parameter value $(-, 0, -, -)$ must be a failure-inducing element. Formally, let $t_{failing}$ be the original failing test case, Δ be the mutation parts, \mathcal{P} be the parameters and their values, then the traditional additional test case generation can be formulated as EQ3.

$$t \leftarrow \text{mutate}(\mathcal{P}, t_{failing}, \Delta) \quad (\text{EQ3})$$

EQ3 indicates that the test case t is generated by mutating the part Δ of the original failing test case $t_{failing}$. Note that the mutated values may have many choices, as long as they are within the scope of \mathcal{P} and different from those in $t_{failing}$. For example, for the original failing test case t_1 (0,0,0,0) in Table 3, let Δ be the second parameter value, then test cases (0, 1, 0, 0) and (0, 2, 0, 0) all satisfy EQ3. We refer to all the test cases that satisfy EQ3 as $\mathcal{T}_{candidate}$, which can be formulated as EQ4.

$$\mathcal{T}_{candidate} = \{ t \mid t \leftarrow \text{mutate}(\mathcal{P}, t_{failing}, \Delta) \} \quad (\text{EQ4})$$

Traditional MFS identification process just selects one test case from $\mathcal{T}_{candidate}$ randomly. However, to adapt the MFS identification process to the new CT framework, this selection should be refined.

Specifically, there are two points to note. First, the additional test case should not contain the currently identified MFS; second, the additional test case is expected to cover as many uncovered schemas as possible. These two goals are similar to CT generation, hence we can directly apply the same selection method on additional test case generation, which can be formulated as EQ5. The same as EQ2, EQ5 excludes the test cases that contain the currently identified MFS from the candidate test cases ($\mathcal{T}_{candidate} - \mathcal{T}_{MFS}$), and selects the additional test case which covers the greatest number of uncovered schemas (Ω).

$$t \leftarrow \text{select}(\mathcal{T}_{candidate} - \mathcal{T}_{MFS}, \Omega, \xi) \quad (\text{EQ5})$$

(3) *Updating uncovered schemas* : After the MFS are identified, some related t -degree schemas, i.e., *MFS themselves*, *super-schemas* and *implicit forbidden schemas*, should be set as covered to enable the termination of the overall CT process. The algorithm that seeks to handle these three types of schemas is listed in Algorithm 1.

Algorithm 1 Changing coverage after identification of MFS

Input: \mathcal{S}_{MFS} ▷ currently identified MFS
 Ω ▷ the schemas that are still uncovered
 \mathcal{T}_{all} ▷ all the possible test cases
 \mathcal{T}_{MFS} ▷ all the test cases that contain the MFS
Output: *void* ▷ do not need any output

```

1: for each  $s \in \mathcal{S}_{MFS}$  do
2:   if  $s$  is  $t$ -degree schema then
3:      $\Omega \leftarrow \Omega \setminus s$ 
4:   end if
5:   for each  $s_p$  is super-schema of  $s$  do
6:     if  $s_p$  is  $t$ -degree schema then
7:        $\Omega \leftarrow \Omega \setminus s_p$ 
8:     end if
9:   end for
10: end for
11: for each  $s \in \Omega$  do
12:   if  $\nexists t \in (\mathcal{T}_{all} - \mathcal{T}_{MFS}), s.t., t.contains(s)$  then
13:      $\Omega \leftarrow \Omega \setminus s$ 
14:   end if
15: end for
```

In this algorithm, we firstly check each MFS (line 1) to see if it is t -degree schema (line 2). We will set those t -degree MFS as covered and remove them from the uncovered schemas set Ω (line 3). This is the first type of schemas – *themselves*. For each t -degree super-schema of these MFS, it will also be removed from the uncovered schemas set (line 5 - 9), as they are the second type of schemas – *super-schemas*. The last type, i.e., *implicit forbidden schemas*, is the toughest one. To remove them, we need to search through each potential schema in the uncovered schemas set (line 11), and check if it is the implicit forbidden schema (line 12). The checking process involves solving a satisfiability problem. Specifically, if we can not find a test case from the set ($\mathcal{T}_{all} - \mathcal{T}_{MFS}$) (excluding those that contain MFS), such that it contains the schema under checking, then we can determine the schema is the implicit forbidden schema and it needs to be removed from uncovered schemas set (line

13). This is because in this case, the schema under checking can only appear in \mathcal{T}_{MFS} , which we will definitely not generate in later iterations. In this paper, a SAT solver will be utilized to do this checking process.

4.3 Advantages of our framework

In view of the above problems listed in Section 3, our new framework has the following advantages:

(1) *Redundant test cases are eliminated such that the overall cost is reduced:*

Two facts of our framework support this improvement: (1) The schemas appeared in the passing test cases is counted towards the overall coverage, so that the test cases generation process has faster convergence rate, which results in generating smaller number of test cases. (2) The forbidden of identified MFS. As a result, test cases which contain these MFS will not appear, as well as those extra-generated test cases used to re-identify these MFS.

(2) *The appearance of multiple MFS in one test case is limited, improving the effectiveness of MFS identification*

This is mainly because we forbidden the appearance of MFS that has been identified. Consequently, with the progress of our approach, the number of remaining MFS decreases one by one. Correspondingly, the probability that multiple MFS appear in one test case will also decrease. Since multiple MFS has a negative effect on MFS identification as discussed in Section 3, the reduction of the appearance of multiple MFS in one test case obviously improve the effectiveness of MFS identification.

(3) *The Masking effects is reduced, hence adequate testing is better satisfied*

As discussed in Section 3, SCT is suffering from masking effects when there are multiple MFS in one failing test case. Since our approach theoretically reduce the probability of the case that multiple MFS appear in one test case, we believe our framework can alleviate the masking effects. In fact, our framework conforms to *tested t-way interaction criterion* because we only update t-way coverage for two types of schemas: (1) t-degree schemas in those passing test cases and (2) t-degree schemas related to MFS. Hence, our framework *Interleaving CT* supports a better adequate testing than SCT.

4.4 Example

With this new framework, when we re-consider the scenario of Section 3, we can get the result listed in Table 4.

This table consists of two main columns, in which the left indicates the generation part while the right column indicates the identification process. We can find that, after identifying the MFS $(-, 0, -, -)$ for t_1 , the following test cases (t_6 to t_{15}) will not contain this schema. Correspondingly, all the 2-degree schemas that are related to this schema, e.g. $(0, 0, -, -)$, $(-, 0, 1, -)$, etc, will also not appear in the following test cases. Additionally, the passing test case t_3 generated in the identification process cover six 2-degree schemas, i.e., $(0, 1, -, -)$, $(0, -, 0, -)$, $(0, -, -, 0)$, $(-, 1, 0, -)$, $(-, 1, -, 0)$, and $(-, -, 0, 0)$ respectively, so that it is not necessary to generate more test cases to cover them. We later found that t_8 failed, which only contain one MFS as expected, and we easily identified it $(-, -, 0, 1)$ with four extra-generated test cases. This schema is 2-degree MFS, which will be forbidden

Table 4: Interleaving CT case study

Generation						Identification					
t_1	0	0	0	0	Fail	t_2^*	1	0	0	0	Fail
						t_3^*	0	1	0	0	Pass
						t_4^*	0	0	1	0	Fail
						t_5^*	0	0	0	1	Fail
						MFS: $(-, 0, -, -)$					
t_6	1	1	1	1	Pass						
t_7	0	2	1	1	Pass						
t_8	0	1	0	1	Fail						
						t_9^*	1	1	0	1	Fail
						t_{10}^*	0	2	0	0	Fail
						t_{11}^*	0	1	1	0	Pass
						t_{12}^*	0	1	0	0	Pass
						MFS: $(-, -, 0, 1)$					
t_{13}	1	2	0	0	Pass						
t_{14}	2	1	0	0	Pass						
t_{15}	2	2	1	1	Pass						

in the following test cases and set to be covered.

Above all, when using the interleaving CT approach, the overall generated test case are 8 less than that of the traditional sequential CT approach in Table 3, and the new approach identified the expected MFS $(-, -, 0, 1)$ which SCT was not able to handle.

5. LIMITATION OF OUR APPROACH

To forbidden identified MFS in the latter generated test cases is efficient for CT, as it will reduce many unnecessary test cases. On the other hand, our framework has strict requirements in accuracy of the identified MFS. This is obviously, as if the schemas identified is not a MFS, latter generated test cases will forbidden a non-MFS schema, which will have two impacts: (1) If this non-MFS schema is the sub-schema of some actual MFS, then the corresponding MFS will never appear and surely we will not detect and identify it. (2) If this non-MFS schema is a sub-schema of some t-degree uncovered schemas, then these schemas will never be covered and a adequate testing will not be reached.

For example, suppose we incorrectly regard $(-, -, 0, -)$ as MFS of the failing test case t_1 in Table 4, then we will never generate test cases contain the real MFS $(-, -, 0, 1)$ (t_8 in Table 4, for example). Hence, $(-, -, 0, 1)$ will be never detected and identified. Additionally, the uncovered 2-degree schemas like $(1, -, 0, -)$, $(-, 2, 0, -)$, $(-, 1, 0, -)$, etc, will never be covered in the latter generated test cases.

To definitely correctly identify the MFS in one failing test case, if possible, however, is not practical due to the cost of testing [27, 10]. This is because for any test case with n parameter values, there are $2^n - 1$ possible schemas which are the candidate MFS. For example, the possible candidate schemas of failing test case $(1, 1, 1)$ are $(1, -, -)$, $(-, 1, -)$, $(-, -, 1)$, $(1, 1, -)$, $(1, -, 1)$, $(-, 1, 1)$ and $(1, 1, 1)$. According to the definition of MFS, we need to individually determine whether these $2^n - 1$ are faulty schemas or not. In fact, even to determine whether a schema is faulty schema or not is not easy, as we must figure out whether all the test cases contain this schema will fail or not. So the complexity to correctly obtain a real MFS is surely exponential. As a result, existing MFS identification approaches actually ob-

tain *approximation* solution through a relatively small size of extra-generated test cases [27, 30, 38, 10, 24, 25, 15].

Based on the insight that incorrectly identifying MFS is inevitable, our approach surely suffer from the two problems discussed previously, but the following measures can help to alleviate such impacts:

1) **Better MFS identification approaches:** Utilizing different approaches is appealing for our framework, as it can improve the robustness of the MFS identification and hence render a more reliable result. Specifically, when a failure is triggered by a test case, instead of only using OFOT, we apply multiple MFS identification approaches (FIC_BS [38], classification tree [35], for example) to identify the MFS. Subsequently, we filters out obtained schemas unless at least two identification approaches produce them as MFS. This simple voting mechanism can improve both precise and robustness of identified MFS. Particularly, the incorrectly MFS can be eliminated if it is obtained by only one identification approach, and MFS that are omitted by one approach can be identified by other approaches. We refer to this mutation version of ICT as *ICT_CB* later.

2) **Tolerating mechanism (Weak forbidden):** Unlike our basic *ICT* framework that rigorously forbidden the identified MFS, we can permit some identified MFS to appear in the latter generated test cases incidentally. This extension can make ICT tolerant of incorrect MFS identification to some extent, as it increases the chance that covering array check them again and re-identify whether they are MFS or not. When to let the identified MFS re-appear in a test case is important, as it affects the efficiency of our framework (If they re-appear too many times, many test cases will fail and MFS identification will repeat again and again, which is inefficiency). In this paper, our implementation of this tolerating mechanism is to inject one identified MFS (The one that has been selected more than twice will be excluded) with a small probability. Later, we refer to this extension version of ICT as *ICT_TL*.

3) More rigorous validation

As we will see in the empirical studies in the next section, although both these two augmentation increase the cost of ICT (i.e., more test cases), they significantly enhance the performance of our framework, especially at the accuracy of MFS identification.

6. SUBJECT PROGRAMS

6.1 Subject programs

The five subject programs used in our experiments are listed in Table 5. Column “Subjects” indicates the specific software. Column “Version” indicates the specific version that is used in the following experiments. Column “LOC” shows the number of source code lines for each software. Column “Faults” presents the fault ID, which is used as the index to fetch the original fault description at each bug tracker for that software. Column “Lan” shows the programming language for each software (For subjects written in more than one programming languages, only the main programming language is shown).

Among these subjects, Tomcat is a web server for java servlet; Hsqldb is a pure-java relational database engine; Gcc is a programming language compiler; Jflex is a lexical analyzer generator; and Tcas is a module of an aircraft collision avoidance system. We select these software as subjects

Table 5: Subject programs

Subjects	Version	LOC	Faults	Lan
Tomcat	7.0.40	296138	#55905	java
Hsqldb	2.0rc8	139425	#981	Java
Gcc	4.7.2	2231564	#55459	c
Jflex	1.4.1	10040	#87	Java
Tcas	v_1	173	#Seed	c

because their behaviours are influenced by various combinations of configuration options or inputs. For example, one component *connector* of Tomcat is influenced by more than 151 attributes [17]. For program Tcas, although with a relatively small size (only 173 lines), it also has 12 parameters with their values ranged from 2 to 10. As a result, the overall input space for Tcas can reach 460800 [32, 19].

As the target of our empirical studies is to compare the ability of fault defection between our approach with traditional ones, we firstly must know these faults and their corresponding MFS in prior, so that we can determine whether the schemas identified by those approaches are accurate or not. For this, we looked through the bug tracker of each software and focused on the bugs which are caused by the interaction of configuration options. Then for each such bug, we derive its MFS by analysing the bug description report and the associated test file which can reproduce the bug. For Tcas, as it does not contain any fault for the original source file, we took a mutation version for that file with injected fault. The mutation was the same as that in [19], which is used as a experimental object for the fault detection studies.

6.2 Specific inputs models

To apply CT on the selected software, we need to firstly model their input parameters. As we discussed before, the whole configuration options is extremely large so that we cannot include all of them in our model in consideration of the experimental time and computing resource. Instead, a moderate small set of these configuration options is selected. It includes the options that cause the specific faults in Table 5, thus the test cases generated by CT can detect these faults. Additional options are also included to create some noise for the MFS identification approach. These options are selected randomly. Details of the specific options and their corresponding values of each software are posted at <http://barbie.uta.edu/data.htm>. A brief overview of the inputs models as well as the corresponding MFS (degree) is shown in Table 6.

Table 6: Inputs model

Subjects	Inputs	MFS
Tomcat	$2^8 \times 3^1 \times 4^1$	1(1) 2(2)
Hsqldb	$2^9 \times 3^2 \times 4^1$	3(2)
Gcc	$2^9 \times 6^1$	3(4)
Jflex	$2^{10} \times 3^2 \times 4^1$	2(1)
Tcas	$2^7 \times 3^2 \times 4^1 \times 10^2$	9(16) 10(8) 11(16) 12(8)

In this table, Column “inputs” depicts the input model for each version of the software, presented in the abbreviated form $\#values^{\#number\ of\ parameters} \times \dots$, e.g., $2^9 \times 3^2 \times 4^1$ indicates the software has 9 parameters that can take 2 values, 2 parameters can take 3 values, and only one parameter that can take 4 values. Column “MFS” shows the degrees

of each MFS and the number of MFS (in the parentheses) with that corresponding degree.

7. EMPIRICAL STUDIES

To evaluate the effectiveness and efficiency of the interleaving CT approach, i.e., *ICT*, we conducted a series of empirical studies on several open-source software subjects. Each of these studies aims at addressing one of the following research questions:

Q1: Does our approach *ICT* perform better than traditional CT process *SCT* at the overall cost and the accuracy of MFS identification?

Q2: Does *ICT* alleviates the three problems proposed in Section 3. Specifically, (1) does *ICT* reduce generating redundant and useless test cases, (2) does *ICT* reduce the appearance of test cases which contains multiple MFS, and (3) does *ICT* reduce the impacts of masking effects?

Q3: Does our approach *ICT* have any advantages over the existed masking effects handling technique — FDA-CIT [36]?

7.1 Comparing ICT with SCT

After preparing the subjects software, next we constructed the experiment to evaluate the efficiency and effectiveness of our approach. To this aim, we need to compare our framework with the traditional sequential CT approach to see if interleaving CT approach has any advantage.

The covering array generating algorithm used in the experiment is AETG [4], as it is the most common one-test-case-one-time generation algorithm. The MFS identifying algorithm is OFOT [27] as discussed before. The constraints handling solver is a java SAT solver – SAT4j [22]. Note that all these three algorithms or techniques can be easily replaced with other similar approaches. For example, we can use other one-test-one-time covering array generation algorithms, like DDA [3], or other MFS identification techniques [38, 30], or other popular SAT solvers [12]. However, to select which algorithms for the three components of combinatorial testing is not the key in this paper, instead, our work focus on the overall CT process.

7.1.1 Study setup

For each software except *Tcas*, a test case is determined to be passing if it ran without any exceptions; otherwise it is regarded as failing. For *Tcas*, as the fault is injected, we determine the result of a test case by separately running it with the original correct version and the mutation version. The test case will be labeled as passing if their results are the same; otherwise, it is deemed as failing.

In this experiment, we focus on three coverage criteria, i.e., 2-way, 3-way and 4-way, respectively. It is known that the generated test cases vary for different runs of AETG algorithm. So to avoid the biases of randomness, we conduct each experiment 30 times and then evaluate the results. (Note that the remaining case studies are also based on 30 repeated experiments.) In other word, for each subject software, we will repeatedly execute traditional approach and our approach 30 times to detect and identify the MFS.

To evaluate the results of the two approaches, one metric is the cost, i.e., the number of test cases that each approach needs. Specifically, the test cases that are generated in the CT generation and MFS identification, respectively, are recorded and compared for these two approaches. Apart

from this, another important metric is the quality of their identified MFS. For this, we used standard metrics: *precision* and *recall*, which are defined as follows:

$$precision = \frac{\#the\ num\ of\ correctly\ identified\ MFS}{\#the\ num\ of\ all\ the\ identified\ schemas}$$

and

$$recall = \frac{\#the\ num\ of\ correctly\ identified\ MFS}{\#the\ num\ of\ all\ the\ real\ MFS}$$

Precision shows the degree of accuracy of the identified schemas when comparing to the real MFS. *Recall* measures how well the real MFS are detected and identified. The combination of them is F-measure, which is

$$F - measure = \frac{2 \times precision \times recall}{precision + recall}$$

7.1.2 Result and discussion

Table 7 presents the results for the number of test cases. In Column ‘Method’, *ict* indicates the interleaving CT approach and *sct* indicates the sequential CT approach. The results of three covering criteria, i.e., 2-way, 3-way, and 4-way are shown in three main columns. In each of them, the number of test cases that are generated in *CT generation* activity (Column ‘Gen’), in *MFS identification* activity (Column ‘Iden’), and the total number of test cases (Column ‘Total’) are listed.

One observation from this table is that the total number of test cases generated by our approach are far less than that of the traditional approach. In fact, except for subject *Tcas*, our approach reduced about dozens of test cases for 2-way coverage, and hundreds of test cases for 3-way and 4-way coverage. For *Tcas*, however, as the MFS are hard to detect (all of them have degrees greater than 9), so both approaches nearly do not trigger errors. Under this condition, both approaches will be transferred to a normal covering array.

The gap of total number of test cases is mainly due to the difference of the number of test cases generated in *MFS identification* activity. In fact, their results in the *CT generation* are almost the same. Even for 4-way coverage criteria which may needs thousands test cases, the gap between them are no more than 10.

For *MFS identification* activity, the interleaving approach only consumed a relatively small amount of test cases when comparing to the sequential approach. What’s more, the interleaving approach obtained almost the same results under the 2-way, 3-way, and 4-way coverage (see the cells in bold). This is as expected, as the MFS identification only happens after a test case fails in our interleaving approach. And after the identification process, the identified MFS will be set as forbidden schemas for the latter generated test cases. As a result, each MFS only needs to be identified once, no matter what the coverage is and how many test cases needed to be generated.

However, this is not the case for the sequential approach. As discussed before, the sequential approach does not identify the MFS at early iteration, so that these MFS can appear in latter test cases. As a result, it needs many more test cases to identify the same MFS, which is a huge waste. Even worse, the more test cases generated in *CT generation* activity, the more test cases are needed in *MFS identification* activity (See the Column ‘Iden’ of sequential approach

Table 7: Comparison of the number of test cases

Subjects	Method	2-way			3-way			4-way		
		Gen	Idea	Total	Gen	Idea	Total	Gen	Idea	Total
Tomcat	ict	12.73	30	42.73	34.97	29.67	64.64	85.57	29.33	114.9
	sct	14.1	105.6	119.7	38.3	256.03	294.33	93.4	578.4	671.8
Hsqldb	ict	14.93	23.2	38.13	45.3	28.0	73.3	118.33	30.4	148.73
	sct	15.63	133.13	148.76	47.37	356.7	404.07	122.7	876.37	999.07
Gcc	ict	14.5	19.33	33.83	47.9	21.33	69.23	102.63	24.67	127.3
	sct	15.27	19.67	34.94	50.97	64.93	115.9	103.87	124.6	228.47
Jflex	ict	15.8	14	29.8	50.47	14	64.47	145.17	14	159.17
	sct	15.83	28.93	44.76	49.63	134.74	184.37	142.1	433.47	575.57
Tcas	ict	109.03	0	109.03	426.8	0.8	427.6	1629.83	3.2	1633.03
	sct	108.67	0	108.67	426.77	1.2	427.97	1633.2	2.8	1636

Table 8: Comparison of the quality of the identified MFS

Subjects	Method	2-way			3-way			4-way		
		Precision	Recall	F-measure	Precision	Recall	F-measure	Precision	Recall	F-measure
Tomcat	ict	1	1	1	0.97	0.96	0.96	0.93	0.91	0.92
	sct	0.73	0.74	0.73	0.75	1	0.86	0.75	1	0.86
Hsqldb	ict	0.52	0.82	0.61	0.56	0.67	0.6	0.69	0.77	0.72
	sct	0.46	0.82	0.57	0.49	0.5	0.49	0.67	0.5	0.57
Gcc	ict	0.47	0.28	0.34	0.46	0.37	0.40	0.58	0.48	0.52
	sct	0.33	0.18	0.23	0.36	0.43	0.39	0.33	0.5	0.40
Jflex	ict	1	1	1	1	1	1	1	1	1
	sct	1	1	1	1	1	1	1	1	1
Tcas	ict	0	0	0	0.03	0.0007	0.001	0.07	0.001	0.0027
	sct	0	0	0	0	0	0	0.03	0.001	0.0026

under 2-way, 3-way and 4-way coverage). This is because the possibility that failures are triggered is increased when there are more test cases without forbidding the appearance of these MFS.

With regarding to the quality of the identified MFS, the comparison of the two approaches are listed in Table 8. Based on this table, we find that there is no apparent gap between them. For example, there are 11 cases under which our approach performed better than traditional one (marked in bold). But among these cases, the maximal gap is 0.27 (2-way for Tomcat), and the average gap is around 0.1, which is trivial.

The reason of the similarity between the quality of these two approaches is that both of them have advantages and disadvantages. Specifically, our approach can reduce the impacts when a test case contains multiple MFS. Our previous study showed that multiple MFS in a test case can reduce the accuracy of the MFS identifying algorithms [30]. As a result, our approach can improve the quality of the identified schemas (details are in the next study). But as a side-effect, if the schemas identified at the early iteration of our approach are not correct, they will significantly impact the following iteration. This is because we will compute the coverage and forbidden schemas based on previous identified MFS. It was the other way around for the traditional approach. It suffers when a test case contains multiple MFS, but correspondingly, previous identified MFS has little influence on the traditional approach. Notes that in the case studies, we do not apply the measures proposed in Section 5, we believe we can obtain a better result if we apply them.

In summary, the answer to **Q1** is: our approach needs much less test cases than traditional sequential CT ap-

proach, and there is no decline in the quality of the identified MFS when comparing with traditional approach.

7.2 Alleviation of the three problems

Section 3 shows three problems that impacts the performance of traditional CT process, which are *redundant test cases generation*, *multiple MFS in one test case* and *masking effects*, respectively. To learn if ICT can alleviate these problems, we re-use the experiment in the first study, i.e., let SCT and ICT generate test cases to identify the MFS in the five program subjects. Then, we respectively investigate the extent to which ICT and SCT that are affected by those impacts.

7.2.1 Study setup

We design three metrics for each of the three problems. First, to measure the *redundant test cases generation*, we gather the number of times that each schema is covered. This metric can directly indicate the redundancy of generated test cases, as it is obvious that if there are too many schemas that are repeated being covered by different test cases, then the CT process is inefficiency as if one schema is covered and tested, it is unnecessary to check them again with other test cases. Second, to measure *multiple MFS in one test case* is simple, we just directly search for each generated test case and examine if it contains more than one MFS. Third, we use the *tested-t-way* coverage criteria [36] to the measure the masking effects. Specifically, we re-compute the coverage of the test cases generated by ICT and SCT by counting all the *t*-degree schemas that is either covered in a passing test case or identified as MFS or faulty schema. For ICT and SCT, the more is the *tested-t-way* coverage,




Figure 3: The redundance of test cases

the more the testing is adequate and hence the less is the masking effects.

7.2.2 Result and discussion

1) Redundant test cases.

Our result is shown in Figure 3. This figure consists of nine sub-figures, one for each subject with specific testing coverage (ranged from 2 - 4 way). For each sub-figure, the x-axis represents the number of times a schema is covered in total, the y-axis represents the number of schemas. For example in the first sub-figure (2-way for Tomcat), two bars with x-coordinate equal to 2 indicates that *ICT* approach have 2.8 schemas in average which are covered twice and *SCT* have 0.4 schemas.

As we discussed previously, the more schemas that are covered low-frequency, the less redundant the generated test cases are. Hence it implies a effective testing if the number of schemas (y-axis) decrease with the increase of the covered times (x-axis). Most of the nine sub-figures indicates that our approach *ict* performs better than *sct*. In fact, for *ict*, the bars decrease rapidly with the increase of x-axis, while *sct*, the trend is more smooth. See subject tomcat with 4-way coverage for example, *ict* have about 80 schemas which are only covered once, about 25 schemas covered twice, latter are decreased to less than 10 schemas when the covered times are bigger than 8. For *sct*, however, for most covered times, it all has about 10 schemas, which indicates a very low effectiveness. The interesting exception is subject Tcas, on which *ict* and *sct* show a similar trend. both *ict* and *sct* does not detect the high-degree MFS in Tcas, and consequently they are both transformed to normal covering arrays.

This result shows that our two modifications of traditional approach, i.e., taking account of the covered schemas by test cases generated in MFS identification and forbidding the appearance of existed MFS to reduce the test cases that are

used to identify the same MFS, are useful, especially when the MFS are detected and identified.

2) Multiple MFS.

The result is shown in Table 9, which lists the number of test cases (on average) which contains more than one MFS.

Table 9: Number of test cases that contain multiple MFS

Subject	Method	2-way	3-way	4-way
Tomcat	ict	0	0	0
	sct	1.13	4.4	10.77
Hsqldb	ict	0	0	0
	sct	0.33	1.87	4.67
Gcc	ict	0.37	0.87	0.47
	sct	0.37	1.6	3.73
Jflex	ict	0	0	0
	sct	0	0	0
Tcas	ict	0	0	0
	sct	0	0	0

From this table, one observation is that our approach *ict* obtained a better result than *sct* at limiting the test cases which contain multiple MFS. For all the subjects except *Gcc*, *ict* eliminated all the test cases which contain multiple MFS. Even if for *Gcc*, the number of test cases which contain MFS is less than approach *sct*. For approach *sct*, however, the result is not as good as *ict*. In fact except subjects *Jflex* and *Tcas*, *sct* suffers from generating test cases which contain multiple MFS. This is why even though *sct* generate much more test cases than *ict*, but it did not obtain a better MFS identification result than *ict*. Two exceptions are subjects are subjects *Jflex* and *Tcas*, on which both *ict* and *sct* did not generate test cases contain multiple MFS. The reason is that *Jflex* only has one MFS (see Table 6) and the MFS of

Table 10: Tested-t-way coverage

Subject	Method	2-way	3-way	4-way
Tomcat	ict	236(100.00%)	1411.17(99.10%)	5459.47(97.49%)
	sct	230.23(97.56%)	1404.7(98.65%)	5550.7(99.12%)
Hsqldb	ict	356.63(99.90%)	2729.57(99.55%)	14051.37(99.40%)
	sct	354.5(99.30%)	2717.63(99.11%)	14065.47(99.50%)
Gcc	ict	249.43(98.98%)	1490.6(97.04%)	5825.43(96.32%)
	sct	251.83(99.93%)	1533.47(99.84%)	6015.8(99.47%)
Jflex	ict	473(100.0%)	4282(100.0%)	26532(100.0%)
	sct	473(100.0%)	4282(100.0%)	26532(100.0%)
Tcas	ict	837(100.0%)	9158(100.0%)	64696(100.0%)
	sct	837(100.0%)	9158(100.0%)	64696(100.0%)

Tcas are all high degree which are hardly detected.

3) Masking effects.

The *tested-t-way* coverage for each approach is shown in Table 10. Specifically, the number of t-degree ($t = 2, 3, 4$) schemas which are *tested* (in the passing test cases or identified as faulty schemas) are gathered, as well as the percentage of the total t-degree schemas (in the parentheses followed). There are several observations can be obtained from this result:

First, the extent to which *sct* and *ict* suffered from masking effects is not severe. Actually, the lowest tested-t-coverage of *ict* is 96.32% (4-way for Gcc) and *sct* is 97.56% (2-way for Tomcat). This is an evidence that combining MFS identification with covering array (either in the sequential way or interleaving way) can make testing more adequate than using covering array alone.

Second, there is not apparent gap between *ict* and *sct* at restricting the masking effects. In our case studies, there are four cases (in bold) that *ict* are better than *sct*, 6 cases that they are equal (for subjects *Jflex* and *Tcas*) and remaining 5 cases that *sct* is better. This gap is trivial. As we discussed before, *ict* forbidden the appearance of identified MFS and can improve the performance at handling masking effects. The results, however, do not show this advantage. One possible explanation for this is that *sct* generate more test cases, of which some are passing test cases that cover more schemas.

Third, if there is only single MFS or MFS is not detected, the teste-t-coverage is 100% for both *ict* and *sct* approaches. This can be observed in subjects *Jflex* and *Tcas*, in which all the t-degree schemas are covered. For single MFS, as MFS identification can accurately identify it out, then all the other t-degree schemas can be determined as irrelevant to the failure. Consequently, the tested-t-coverage is satisfied. For the case that MFS cannot be detected (MFS with high degrees), traditional t-way covering arrays are enough to obtain adequate testing, as all the test cases will pass if there is no MFS detected.

To summary, the answer to **EQ 2** is we can learn that our approach can alleviate the three problem discussed in Section 3, especially on reducing the redundancy of test cases and limiting the test cases with multiple MFS. Additionally, both *ict* and *sct* have a good performance on reducing the masking effects.

7.3 Comparison with FDA-CIT

FDA-CIT [36] is a feedback framework that can augment the traditional covering array to iteratively identify the MF-

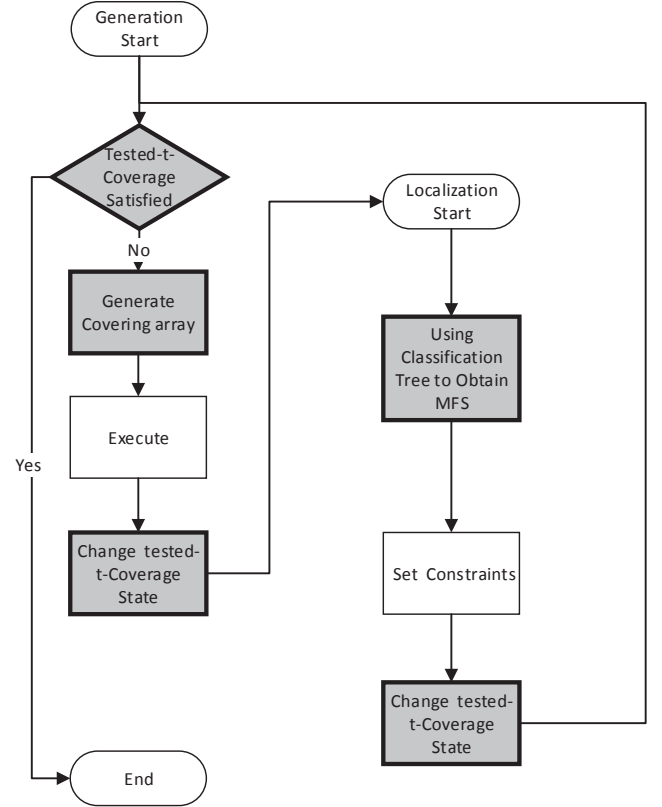


Figure 4: The Framework of FDA-CIT

S, and can handle the masking effects. The overall process can be illustrated in Figure 4. Specifically, it will first generate a t-way covering array and executed all the test cases in it. After that it will utilize the classification tree method to identify the MFS. Then it will forbidden the identified MFS to be appeared and compute the tested-t-coverage. If the tested-t-coverage is not satisfied, it will repeat the previous process, i.e., generating additional test cases and identifying MFS. Like our *ict* approach, FDA-CIT is also an adaptive approach which iteratively generate test cases and identify the MFS. Besides these commonalities, there are several important difference between our approach and FDA-CIT (marked with darker background in Figure 4):

First, the **granularity** of adaptation. Instead of handling one test case one time as *ict*, FDA-CIT tries to generate a batch of test cases at each iteration (A complete covering array will be generated at first iteration, and more test cases will be supplemented to cover those t-degree schemas which are masked at latter iterations). To generate a batch of test cases may improve the degree of parallelism of testing, but this coarser granularity may also introduce some problems, e.g., some test cases generated at one iteration may fail with the same MFS, which is a potential waste because it is better to use one failing test case to reveal one particular MFS.

Second, the **MFS identification** approach is different. FDA-CIT uses classification tree on existed executed test cases to characterize the MFS. Different from our OFOT approach, this post-analysis technique does not need additional test cases, but as a side effect it cannot precisely find the MFS. What's worse, the effectiveness of this post-analysis approach depends greatly on how the covering array is, e.g., if there is large number of failing tests, and small size of the test suite, there is little information to exclude the particular MFS [38].

Third, the **coverage criteria** is not the same. FDA-CIT directly use the tested-t-way coverage to guide their process. This support a better adequate testing and reduce the impacts of masking effects. As we will see later in our experiments, however, the incorrectly MFS identification may negatively affected the achievement of this type of coverage.

Next, we will design experiments to measure the performance of FDA-CIT when handling the MFS in those subjects in Table 5.

7.3.1 Study setup

The design of this case study is similar as previous two. For each subject in Table 5, we apply FDA-CIT to generate test cases and identify the MFS in it. After that, we gather the overall test cases generated (FDA-CIT does not need to additional test cases to identify the MFS), MFS identification results(including re-call, precise and f-measure), and other two metrics, i.e., covered times of schemas, the t-tested-coverage. Note that in this study, we do not gather the test cases which contain multiple MFS, as the post-analysis classification tree method are not affected by this factor. The same as previous experiments, we will repeated executing each experiment 30 times for different coverage (2, 3, and 4 way), and then gather and analyse the average data.

7.3.2 Result and discussion

1)**Overall number of test cases** The overall number of test cases generated by FDA-CIT for each subject is shown

in Table 11. To better evaluate the performance of FDA-CIT, we list the gaps between FDA-CIT with *ict* and *sct* respectively in the parentheses (the first one is *ict*, latter one is *sct*). The computation of this gap between two approaches result is by .

From this table, one observation is that FDA-CIT are better than *sct* at all these cases. Combining the results of previous studies for *sct* and *ict*, we can conclude that the *sct* is most inefficient approach at test cases generation. Second for approach *ict* and FDA-CIT, there are ups and downs on both sides. In detail, FDA-CIT needs less test cases at lower coverage (Except *tcas*, fda-cit reduce about 50% number of test cases when comparing with *ict* for 2-way coverage), while *ict* performs better at higher coverage (Except *tcas* and *gcc* which two approaches have similar results, *fda-cit* increases about 10% number of test cases).

This result is reasonable. First, FDA-CIT does not need additional test cases to identify the MFS, which will reduce some cost when comparing with *ict*, especially when the coverage is low (For low coverage, the test cases generated by *ict* in the MFS identification stage accounts for a considerable proportion of the overall test cases). One the other hand, as we said before, the coarse grained generation will make fda-cit generate some unnecessary test cases.

2)**the F-measure of MFS identification** The results of the quality of MFS identification by FDA-CIT is listed in Table 12. Then same as previous metric, we also list the comparison between fda-cit with *ict* and *sct* are also attached.

This table shows a very clear gap between. In fact, High degree is not well.

3)**Redundant test cases.**

The result is listed in Figure 5.

4)**Masking effects.**

The results is listed in Table 13.

7.4 Threats to validity

There are several threats to validity in our empirical studies. First, our experiments are based on only 5 open-source software. More subject programs are desired to make the results more general. In fact, we plan to conduct comprehensive experiments on the programs with parameters and MFS under control, such that the conclusion of our experiment can reduce the impact caused by specific input space and specific degree or location of the MFS.

Second, there have been many more generation algorithms and MFS identification algorithms. In our empirical studies, we just used AETG [4] as the test case generation strategy, and OFOT [27] as the MFS identification strategy. As different generation and identification algorithms may affect the performance our proposed CT framework, especially on the number of test cases, some studies for different test case generation and MFS identification approaches are desired.

8. RELATED WORKS

Combinatorial testing has been widely applied in practice [18], especially on domains like configuration testing [35, 9, 13] and software inputs testing [4, 1, 14]. A recent survey [28] comprehensively studied existing works in CT and classified those works into eight categories according to the testing procedure. Based on which, we can learn that test cases generation and MFS identification are two most important parts in CT studies.

Table 11: Number of test cases generated by FDA-CIT

	2-way	3-way	4-way
Tomcat	27.93 (52.97%, 328.52%)	66.57 (-2.89%, 342.16%)	201.1 (-42.86%, 234.06%)
HSQldb	28.67 (33.02%, 416.74%)	86 (-14.77%, 367.60%)	179.77 (-17.26%, 454.24%)
GCC	20.63 (63.96%, 69.34%)	63.07 (9.77%, 83.77%)	125.1 (1.76%, 82.63%)
Jflex	19.57 (52.30%, 128.76%)	65.4 (-1.42%, 181.91%)	181.3 (-12.21%, 217.47%)
Tcas	19.57 (0.89%, 0.56%)	65.4 (2.76%, 2.84%)	181.3 (4.62%, 4.81%)

Table 12: The F-measure of MFS identification for FDA-CIT

	2-way	3-way	4-way
Tomcat	0.26 (73.56%, 46.56%)	0.31 (64.62%, 54.62%)	0.33 (58.67%, 52.67%)
HSQldb	0.33 (27.80%, 24.02%)	0.26 (33.94%, 22.70%)	0.31 (40.77%, 25.91%)
GCC	0.08 (25.83%, 14.83%)	0.44 (-4.44%, -5.44%)	0.47 (5.45%, -6.55%)
Jflex	1 (0.00%, 0.00%)	1 (0.00%, 0.00%)	1 (0.00%, 0.00%)
Tcas	0 (0.00%, 0.00%)	0 (0.10%, 0.00%)	0 (0.27%, 0.26%)

Although CT has been proven to be effective at detecting and identifying the interaction failures in SUT, however, to directly apply them in practice can be inefficient and some times even does not work at all. Some problems, e.g., constraints of parameters values in SUT [6, 8], masking effects of multiple failures [11, 36], dynamic requirement for the strength of covering array [13], will bring many troubles to CT process. To overcome these problems, some works try to make CT more adaptive and flexible.

JieLi [24] augmented the MFS identifying algorithm by selecting one previous passing test case for comparison, such that it can reduce some extra test cases when compared to another efficient MFS identifying algorithm [38].

S.Fouché et al., [13] introduced the notion of incremental covering. Different from traditional covering array, it does not need a fixed strength to guide the generation; instead, it can dynamically generate high-way covering array based on existing low-way covering array, which can support a flexible tradeoff between covering array strength and testing resources. Cohen [6, 8] studied the impacts of constraints for CT, and proposed an SAT-based approach that can handle those constraints. Bryce and Colbourn [2] proposed an one-test-case-one-time greedy technique to not only generate test cases to cover all the t -degree interactions, but also prioritize them according their importance. E. Dumlu et al., [11] developed a feedback driven combinatorial testing approach that can assist traditional covering in avoiding the masking effects between multiple failures. Yilmaz [36] extended that work by refining the MFS diagnosing method. Specifically, this feedback driven approach firstly generate a t -way covering array, and after executing them, the MFS will be identified by utilizing a classification tree method. It then forbidden these MFS and generate additional test cases to cover the interactions that are masked by the MFS. This process continues until all the interactions are covered. Additionally, Nie [29] constructed an adaptive combinatorial testing framework, which can dynamically adjust the inputs model, strength t of covering array, and generation strategy during CT process.

Our work differs from them mainly in that we proposed a highly interactive framework for test case generation and MFS identification. Specifically, we do not generate a complete t -way covering array at first; instead, when a failure is triggered by a test case, we immediately terminate test

cases generation and turn to MFS identification. After the MFS is identified, the coverage will be updated and the test cases generation process continues.

9. CONCLUSIONS

Combinatorial testing is an effective testing technique at detecting and diagnosis the failure-inducing interactions in the SUT. Traditional CT separately studies test cases generation and MFS identification. In this paper, we proposed a new CT framework, i.e., *interleaving CT*, that integrates these two important stages, which allows for both generation and identification better share each other's information. As a result, interleaving CT approach can provide a more efficient testing than traditional sequential CT.

Empirical studies were conducted on five open-source software subjects. The results show that with our new CT framework, there is a significant reduction on the number of generated test cases when compared to the traditional sequential CT approach, while there is no decline in the quality of the identified MFS. Further, when comparing with the ELA [26, 25], our approach also performed better, especially with fewer test cases.

As a future work, we need to extend our interleaving CT approach with more test case generation and MFS identification algorithms, to see the extent on which our new CT framework can enhance those different CT-based algorithms. Another interesting work is to combine interleaving CT approach with white-box testing techniques, so that it can provide more useful information for developers to debug the system.

10. REFERENCES

- [1] M. N. Borazjany, L. Yu, Y. Lei, R. Kacker, and R. Kuhn. Combinatorial testing of acts: A case study. In *Software Testing, Verification and Validation (ICST), 2012 IEEE Fifth International Conference on*, pages 591–600. IEEE, 2012.
- [2] R. C. Bryce and C. J. Colbourn. Prioritized interaction testing for pair-wise coverage with seeding and constraints. *Information and Software Technology*, 48(10):960–970, 2006.
- [3] R. C. Bryce and C. J. Colbourn. The density algorithm for pairwise interaction testing. *Software*

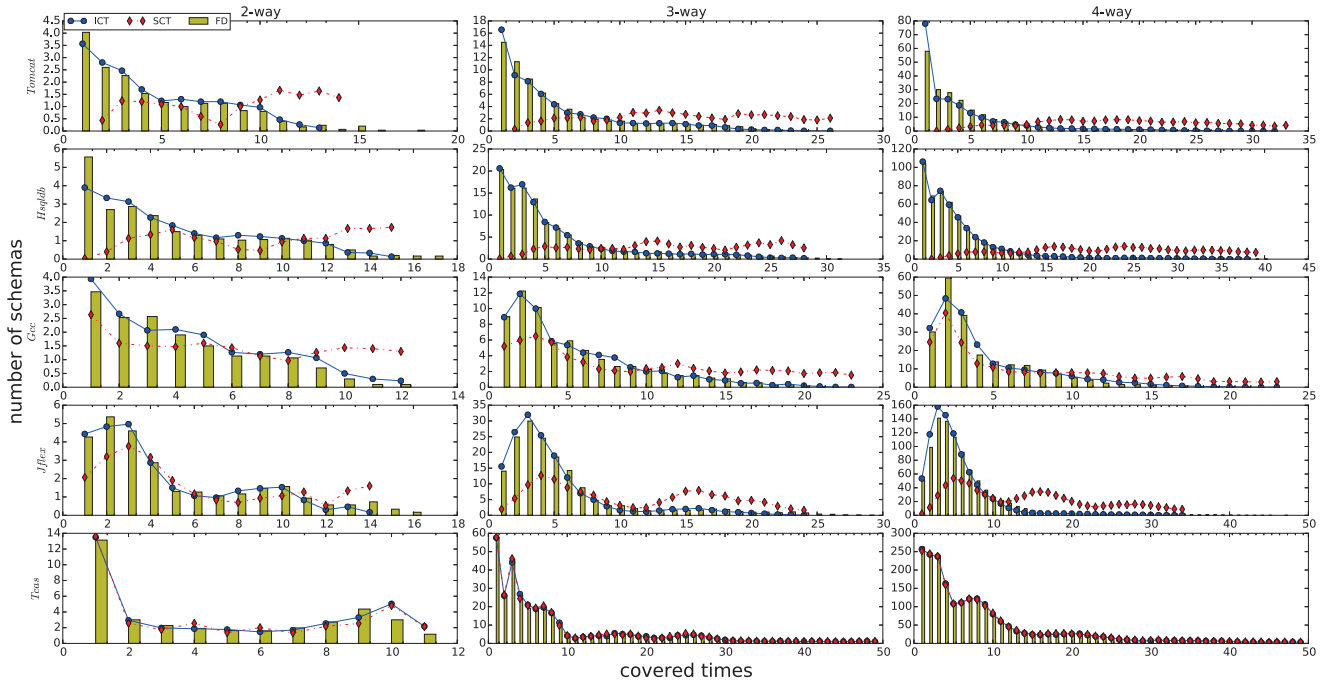


Figure 5: Result of ex2

Table 13: The tested-t-way coverage for FDA-CIT

	2-way	3-way	4-way
Tomcat	208.2 (11.78%, 9.34%)	1387.3 (1.68%, 1.22%)	5378 (1.46%, 3.08%)
HSQLDB	342.1 (4.07%, 3.47%)	2732.43 (-0.11%, -0.54%)	14093.17 (-0.30%, -0.20%)
GCC	250.37 (-0.37%, 0.58%)	1532.03 (-2.70%, 0.09%)	6016.97 (-3.17%, -0.02%)
Jflex	473 (0.00%, 0.00%)	4282 (0.00%, 0.00%)	26532 (0.00%, 0.00%)
Tcas	837 (0.00%, 0.00%)	9158 (0.00%, 0.00%)	64696 (0.00%, 0.00%)

Testing, Verification and Reliability, 17(3):159–182, 2007.

- [4] D. M. Cohen, S. R. Dalal, M. L. Fredman, and G. C. Patton. The aetg system: An approach to testing based on combinatorial design. *Software Engineering, IEEE Transactions on*, 23(7):437–444, 1997.
- [5] M. B. Cohen, C. J. Colbourn, and A. C. Ling. Augmenting simulated annealing to build interaction test suites. In *Software Reliability Engineering, 2003. ISSRE 2003. 14th International Symposium on*, pages 394–405. IEEE, 2003.
- [6] M. B. Cohen, M. B. Dwyer, and J. Shi. Exploiting constraint solving history to construct interaction test suites. In *Testing: Academic and Industrial Conference Practice and Research Techniques-MUTATION, 2007.*, pages 121–132. IEEE, 2007.
- [7] M. B. Cohen, M. B. Dwyer, and J. Shi. Interaction testing of highly-configurable systems in the presence of constraints. In *Proceedings of the 2007 international symposium on Software testing and analysis*, pages 129–139. ACM, 2007.
- [8] M. B. Cohen, M. B. Dwyer, and J. Shi. Constructing interaction test suites for highly-configurable systems in the presence of constraints: A greedy approach. *Software Engineering, IEEE Transactions on*, 34(5):633–650, 2008.
- [9] M. B. Cohen, J. Snyder, and G. Rothermel. Testing across configurations: implications for combinatorial testing. *ACM SIGSOFT Software Engineering Notes*, 31(6):1–9, 2006.
- [10] C. J. Colbourn and D. W. McClary. Locating and detecting arrays for interaction faults. *Journal of combinatorial optimization*, 15(1):17–48, 2008.
- [11] E. Dumlu, C. Yilmaz, M. B. Cohen, and A. Porter. Feedback driven adaptive combinatorial testing. In *Proceedings of the 2011 International Symposium on Software Testing and Analysis*, pages 243–253. ACM, 2011.
- [12] N. Eén and N. Sörensson. An extensible sat-solver. In *Theory and applications of satisfiability testing*, pages 502–518. Springer, 2004.
- [13] S. Fouché, M. B. Cohen, and A. Porter. Incremental covering array failure characterization in large configuration spaces. In *Proceedings of the eighteenth international symposium on Software testing and analysis*, pages 177–188. ACM, 2009.
- [14] B. Garn and D. E. Simos. Eris: A tool for combinatorial testing of the linux system call interface. In *Software Testing, Verification and Validation Workshops (ICSTW), 2014 IEEE Seventh*

- International Conference on*, pages 58–67. IEEE, 2014.
- [15] L. S. G. Ghandehari, Y. Lei, T. Xie, R. Kuhn, and R. Kacker. Identifying failure-inducing combinations in a combinatorial test set. In *Software Testing, Verification and Validation (ICST), 2012 IEEE Fifth International Conference on*, pages 370–379. IEEE, 2012.
 - [16] D. Jin, X. Qu, M. B. Cohen, and B. Robinson. Configurations everywhere: Implications for testing and debugging in practice. In *Companion Proceedings of the 36th International Conference on Software Engineering*, pages 215–224. ACM, 2014.
 - [17] K. Kolinko. The HTTP Connector. <http://tomcat.apache.org/tomcat-7.0-doc/config/http.html>, 2014. [Online; accessed 3-Nov-2014].
 - [18] D. R. Kuhn, R. N. Kacker, and Y. Lei. Practical combinatorial testing. *NIST Special Publication*, 800:142, 2010.
 - [19] D. R. Kuhn and V. Okun. Pseudo-exhaustive testing for software. In *Software Engineering Workshop, 2006. SEW’06. 30th Annual IEEE/NASA*, pages 153–158. IEEE, 2006.
 - [20] D. R. Kuhn and M. J. Reilly. An investigation of the applicability of design of experiments to software testing. In *Software Engineering Workshop, 2002. Proceedings. 27th Annual NASA Goddard/IEEE*, pages 91–95. IEEE, 2002.
 - [21] D. R. Kuhn, D. R. Wallace, and J. AM Gallo. Software fault interactions and implications for software testing. *Software Engineering, IEEE Transactions on*, 30(6):418–421, 2004.
 - [22] D. Le Berre, A. Parrain, et al. The sat4j library, release 2.2, system description. *Journal on Satisfiability, Boolean Modeling and Computation*, 7:59–64, 2010.
 - [23] Y. Lei, R. Kacker, D. R. Kuhn, V. Okun, and J. Lawrence. Ipog/ipog-d: efficient test generation for multi-way combinatorial testing. *Software Testing, Verification and Reliability*, 18(3):125–148, 2008.
 - [24] J. Li, C. Nie, and Y. Lei. Improved delta debugging based on combinatorial testing. In *Quality Software (QSIC), 2012 12th International Conference on*, pages 102–105. IEEE, 2012.
 - [25] C. Martínez, L. Moura, D. Panario, and B. Stevens. Algorithms to locate errors using covering arrays. In *LATIN 2008: Theoretical Informatics*, pages 504–519. Springer, 2008.
 - [26] C. Martínez, L. Moura, D. Panario, and B. Stevens. Locating errors using elas, covering arrays, and adaptive testing algorithms. *SIAM Journal on Discrete Mathematics*, 23(4):1776–1799, 2009.
 - [27] C. Nie and H. Leung. The minimal failure-causing schema of combinatorial testing. *ACM Transactions on Software Engineering and Methodology (TOSEM)*, 20(4):15, 2011.
 - [28] C. Nie and H. Leung. A survey of combinatorial testing. *ACM Computing Surveys (CSUR)*, 43(2):11, 2011.
 - [29] C. Nie, H. Leung, and K.-Y. Cai. Adaptive combinatorial testing. In *Quality Software (QSIC), 2013 13th International Conference on*, pages 284–287. IEEE, 2013.
 - [30] X. Niu, C. Nie, Y. Lei, and A. T. Chan. Identifying failure-inducing combinations using tuple relationship. In *Software Testing, Verification and Validation Workshops (ICSTW), 2013 IEEE Sixth International Conference on*, pages 271–280. IEEE, 2013.
 - [31] X. Qu, M. B. Cohen, and G. Rothermel. Configuration-aware regression testing: an empirical study of sampling and prioritization. In *Proceedings of the 2008 international symposium on Software testing and analysis*, pages 75–86. ACM, 2008.
 - [32] K. Shakya, T. Xie, N. Li, Y. Lei, R. Kacker, and R. Kuhn. Isolating failure-inducing combinations in combinatorial testing using test augmentation and classification. In *Software Testing, Verification and Validation (ICST), 2012 IEEE Fifth International Conference on*, pages 620–623. IEEE, 2012.
 - [33] C. Song, A. Porter, and J. S. Foster. itree: Efficiently discovering high-coverage configurations using interaction trees. In *Proceedings of the 2012 International Conference on Software Engineering*, pages 903–913. IEEE Press, 2012.
 - [34] Z. Wang, B. Xu, L. Chen, and L. Xu. Adaptive interaction fault location based on combinatorial testing. In *Quality Software (QSIC), 2010 10th International Conference on*, pages 495–502. IEEE, 2010.
 - [35] C. Yilmaz, M. B. Cohen, and A. A. Porter. Covering arrays for efficient fault characterization in complex configuration spaces. *Software Engineering, IEEE Transactions on*, 32(1):20–34, 2006.
 - [36] C. Yilmaz, E. Dumlu, M. Cohen, and A. Porter. Reducing masking effects in combinatorial interaction testing: A feedback driven adaptive approach. *Software Engineering, IEEE Transactions on*, 40(1):43–66, Jan 2014.
 - [37] S. Yoo, M. Harman, and D. Clark. Fault localization prioritization: Comparing information-theoretic and coverage-based approaches. *ACM Transactions on Software Engineering and Methodology (TOSEM)*, 22(3):19, 2013.
 - [38] Z. Zhang and J. Zhang. Characterizing failure-causing parameter interactions by adaptive testing. In *Proceedings of the 2011 International Symposium on Software Testing and Analysis*, pages 331–341. ACM, 2011.